



Ph.D. defense

ALICE

Helga

Introduction

Problem
space

Subspace of
instances

Feature space

Algorithm
space

Performance
space

Footprints in
instance space

Preference set

Preference
learning

Conclusions

ALICE

Analysis & Learning Iterative Consecutive Executions

Helga Ingimundardóttir

University of Iceland

June 30, 2016



Introduction

ALICE

Helga

Introduction

Problem
space

Subspace of
instances

Feature space

Algorithm
space

Performance
space

Footprints in
instance space

Preference set

Preference
learning

Conclusions

Motivation:

- ★ The general goal is to train optimisation algorithms, for an arbitrary problem domain, using data.

Contribution:

- ★ The main contribution of this thesis is towards a better understanding of how this training data should be constructed.

Introduction

ALICE

Helga

Introduction

Problem
space

Subspace of
instances

Feature space

Algorithm
space

Performance
space

Footprints in
instance space

Preference set

Preference
learning

Conclusions

Motivation:

- ★ The general goal is to train optimisation algorithms, for an arbitrary problem domain, using data.

Contribution:

- ★ The main contribution of this thesis is towards a better understanding of how this training data should be constructed.

Framework for Algorithm Learning

ALICE

Helga

Introduction

Problem space

Subspace of instances

Feature space

Algorithm space

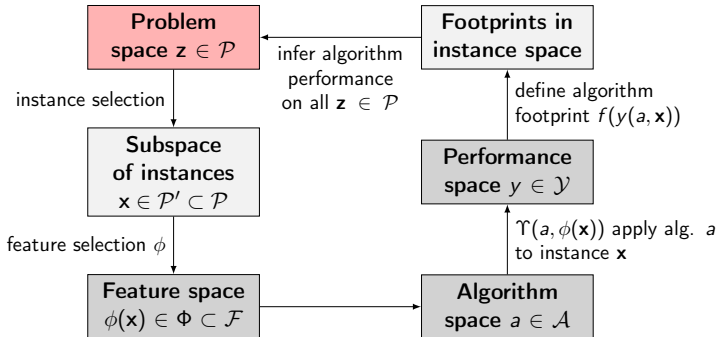
Performance space

Footprints in instance space

Preference set

Preference learning

Conclusions



Framework for Algorithm Learning

ALICE

Helga

Introduction

Problem
space

Subspace of
instances

Feature space

Algorithm
space

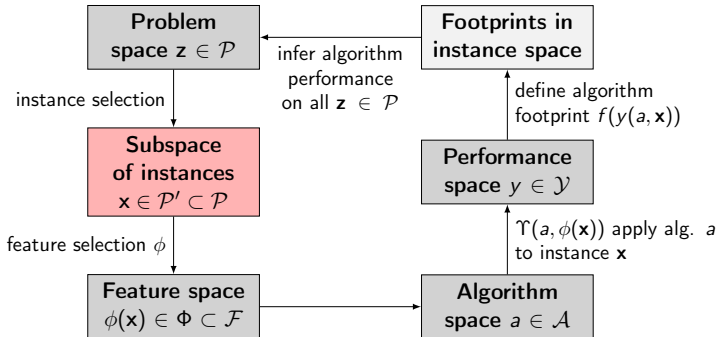
Performance
space

Footprints in
instance space

Preference set

Preference
learning

Conclusions



Framework for Algorithm Learning

ALICE

Helga

Introduction

Problem space

Subspace of instances

Feature space

Algorithm space


Performance space

Footprints in instance space

Preference set

Preference learning

Conclusions



Framework for Algorithm Learning

ALICE

Helga

Introduction

Problem
space

Subspace of
instances

Feature space

Algorithm
space

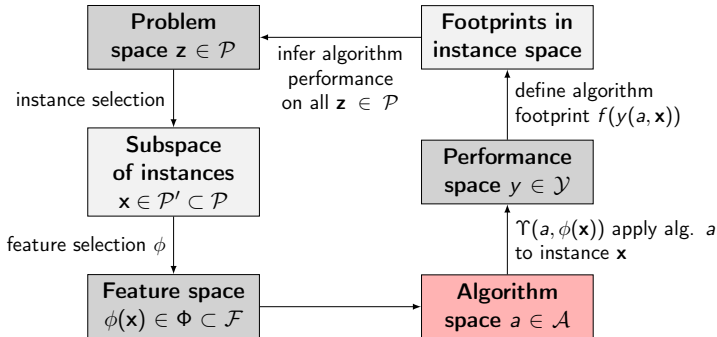
Performance
space

Footprints in
instance space

Preference set

Preference
learning

Conclusions



Framework for Algorithm Learning

ALICE

Helga

Introduction

Problem
space

Subspace of
instances

Feature space

Algorithm
space

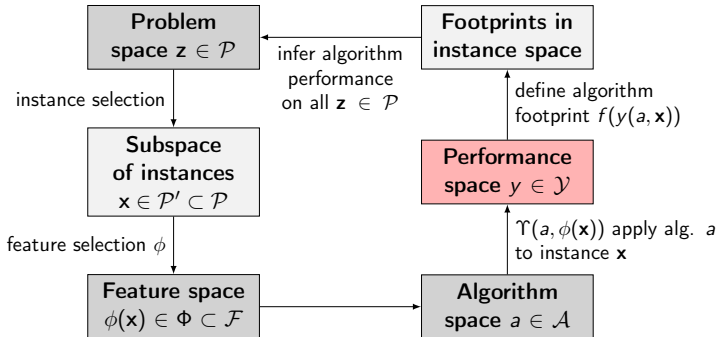
Performance
space

Footprints in
instance space

Preference set

Preference
learning

Conclusions



Framework for Algorithm Learning

ALICE

Helga

Introduction

Problem
space

Subspace of
instances

Feature space

Algorithm
space

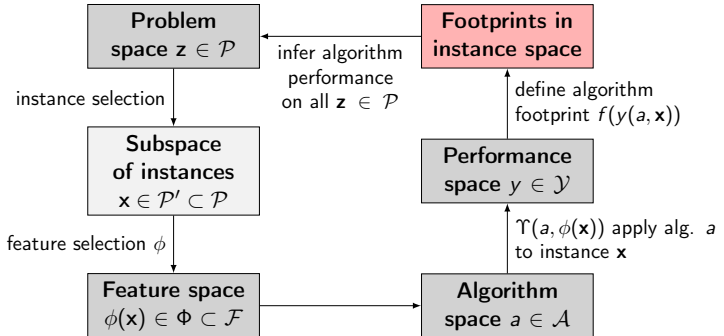
Performance
space

Footprints in
instance space

Preference set

Preference
learning

Conclusions



Framework for Algorithm Learning

ALICE

Helga

Introduction

Problem space

Subspace of instances

Feature space

Algorithm space


Performance space

Footprints in instance space

Preference set

Preference learning

Conclusions



Framework for Algorithm Learning

ALICE

Helga

Introduction

Problem space

Subspace of instances

Feature space

Algorithm space

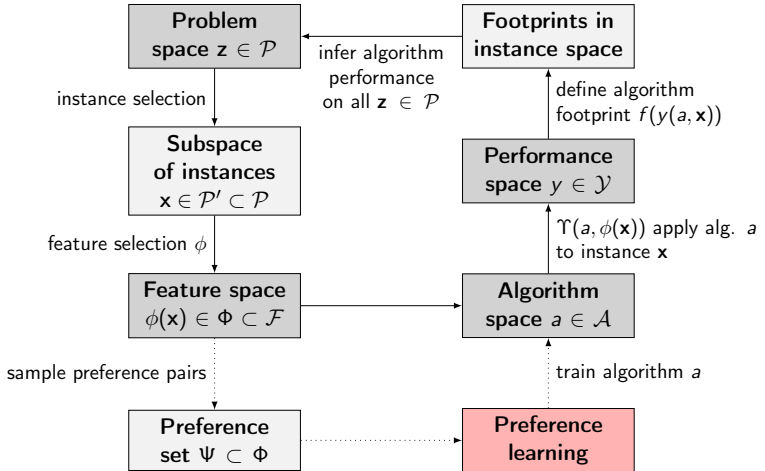
Performance space

Footprints in instance space

Preference set

Preference learning

Conclusions



Framework for Algorithm Learning

ALICE

Helga

Introduction

Problem
space

Subspace of
instances

Feature space

Algorithm
space

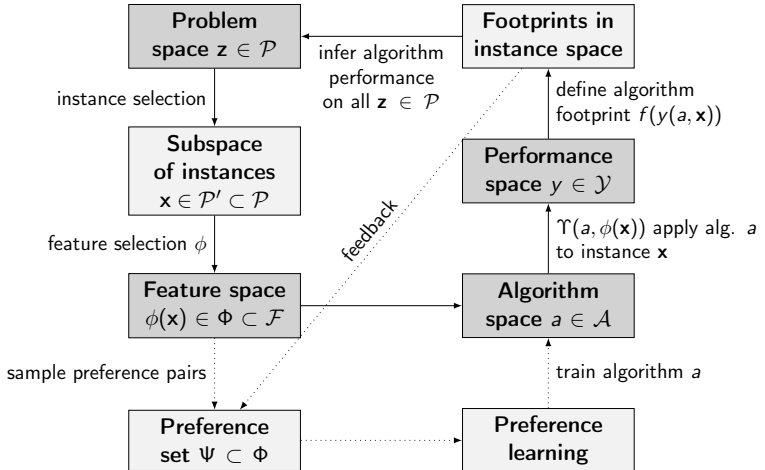
Performance
space

Footprints in
instance space

Preference set

Preference
learning

Conclusions



Framework for Algorithm Learning

ALICE

Helga

Introduction

Problem space

Subspace of instances

Feature space

Algorithm space


Performance space

Footprints in instance space

Preference set

Preference learning

Conclusions





Job Shop Scheduling (JSP)

ALICE

Helga

Introduction

**Problem
space**

Subspace of
instances

Feature space

Algorithm
space

Performance
space

Footprints in
instance space

Preference set

Preference
learning

Conclusions

Simple job-shop is where n jobs are scheduled on a set of m machines, subject to constraints:

- ★ each job must follow a predefined machine order,
- ★ that a machine can handle at most one job at a time.

Objective: schedule the jobs so as to minimise the maximum completion time, i.e., makespan, C_{\max} .

Dispatching rules

ALICE

Helga

Introduction

Problem
space

Subspace of
instances

Feature space

Algorithm
space

Performance
space

Footprints in
instance space

Preference set

Preference
learning

Conclusions

Dispatching rules (DR) are consecutive executions found by:

- ★ Starting with an empty schedule and adding on one operation at a time.
- ★ When a machine is free the DR inspects the available jobs and selects the one with the **highest priority**.
- ★ Complete schedule consists of $\ell = n \cdot m$ sequential dispatches.
- ★ At each dispatch k , **features** $\phi(k)$ for the temporal schedule are calculated.



Mad Hatter tea-party (definition)

ALICE

Helga

Introduction

Problem
space

Subspace of
instances

Feature space

Algorithm
space

Performance
space

Footprints in
instance space

Preference set

Preference
learning

Conclusions

The attending guests: They all have to:

J_1) Alice

M_1) have wine or pour tea

J_2) March Hare

M_2) spread butter

J_3) Dormouse

M_3) get a haircut

J_4) Mad Hatter.

M_4) check the time of the broken watch

M_5) say what they mean.

This can be considered as is a typical 4×5 job-shop, where:

- ★ our guests are the jobs
- ★ their tasks are the machines
- ★ objective is to minimise C_{\max} (when Alice can leave).

Mad Hatter tea-party (action states)

ALICE

Helga

Introduction

Problem space

Subspace of instances

Feature space

Algorithm space

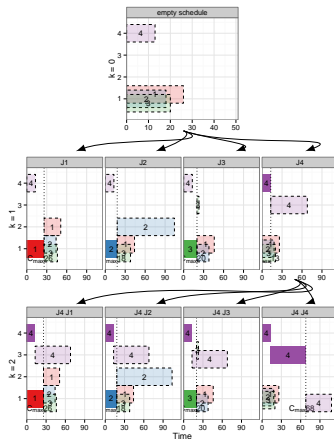


Performance space

Footprints in instance space

Preference set

Preference learning

Conclusions



Mad Hatter tea-party (action states)

ALICE

Helga

Introduction

Problem space

Subspace of instances

Feature space

Algorithm space

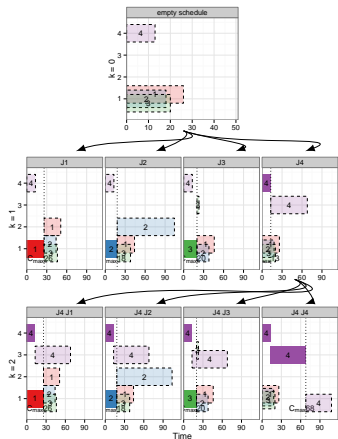

Performance space

Footprints in instance space

Preference set

Preference learning

Conclusions



Mad Hatter tea-party (action states)

ALICE

Helga

Introduction

Problem space

Subspace of instances

Feature space

Algorithm space




Performance space

Footprints in instance space

Preference set

Preference learning

Conclusions



Mad Hatter tea-party (SDRs)

ALICE

Helga

Introduction

Problem
space

Subspace of
instances

Feature space

Algorithm
space


Performance
space

Footprints in
instance space

Preference set

Preference
learning

Conclusions



Framework for Algorithm Learning

ALICE

Helga

Introduction

Problem
space

Subspace of
instances

Feature space

Algorithm
space

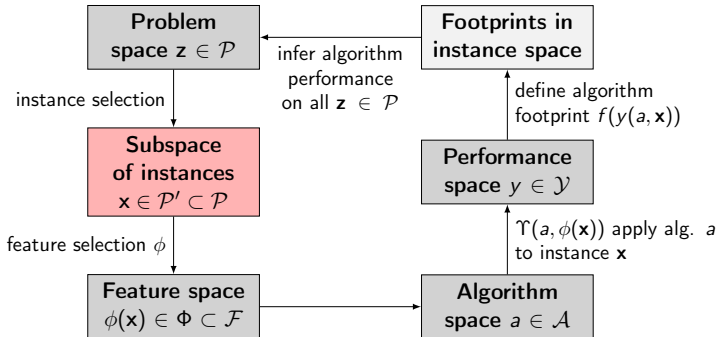
Performance
space

Footprints in
instance space

Preference set

Preference
learning

Conclusions





Problem instance generators

ALICE

Helga

Introduction

Problem space

Subspace of instances

Feature space

Algorithm space

Performance space

Footprints in instance space

Preference set

Preference learning

Conclusions

	name	size ($n \times m$)	N_{train}	N_{test}	note
JSP	$\mathcal{P}_{j.\text{rnd}}^{6 \times 5}$	6×5	500	500	random
	$\mathcal{P}_{j.\text{rndn}}^{6 \times 5}$	6×5	500	500	random-narrow
	$\mathcal{P}_{j.\text{rnd}, J_1}^{6 \times 5}$	6×5	500	500	random with job variation
	$\mathcal{P}_{j.\text{rnd}, M_1}^{6 \times 5}$	6×5	500	500	random with machine variation
	$\mathcal{P}_{j.\text{rnd}}^{10 \times 10}$	10×10	300	200	random
	$\mathcal{P}_{j.\text{rndn}}^{10 \times 10}$	10×10	300	200	random-narrow
	$\mathcal{P}_{j.\text{rnd}, J_1}^{10 \times 10}$	10×10	300	200	random with job variation
	$\mathcal{P}_{j.\text{rnd}, M_1}^{10 \times 10}$	10×10	300	200	random with machine variation
	$\mathcal{P}_{\text{JSP.ORLIB}}$	various	–	82	various
FSP	$\mathcal{P}_{f.\text{rnd}}^{6 \times 5}$	6×5	500	500	random
	$\mathcal{P}_{f.\text{rndn}}^{6 \times 5}$	6×5	500	500	random-narrow
	$\mathcal{P}_{f.\text{jc}}^{6 \times 5}$	6×5	500	500	job-correlated
	$\mathcal{P}_{f.\text{mc}}^{6 \times 5}$	6×5	500	500	machine-correlated
	$\mathcal{P}_{f.\text{mxc}}^{6 \times 5}$	6×5	500	500	mixed-correlation
	$\mathcal{P}_{f.\text{rnd}}^{10 \times 10}$	10×10	300	200	random
	$\mathcal{P}_{\text{FPS.ORLIB}}$	various	–	31	various

Framework for Algorithm Learning

ALICE

Helga

Introduction

Problem
space

Subspace of
instances

Feature space

Algorithm
space

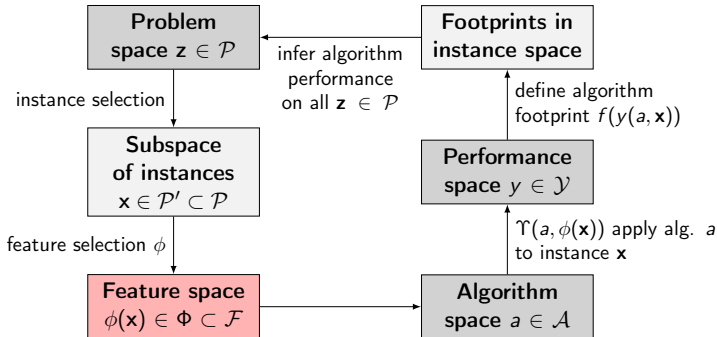
Performance
space

Footprints in
instance space

Preference set

Preference
learning

Conclusions



Features for JSP

ALICE

Helga

Introduction

Problem
space

Subspace of
instances

Feature space

Algorithm
space

Performance
space

Footprints in
instance space

Preference set

Preference
learning

Conclusions

job	ϕ_1	job processing time
	ϕ_2	job start-time
	ϕ_3	job end-time
	ϕ_4	job arrival time
	ϕ_5	time job had to wait
	ϕ_6	total processing time for job
	ϕ_7	total work remaining for job
	ϕ_8	number of assigned operations for job
machine	ϕ_9	when machine is next free
	ϕ_{10}	total processing time for machine
	ϕ_{11}	total work remaining for machine
	ϕ_{12}	number of assigned operations for machine
	ϕ_{13}	change in idle time by assignment
	ϕ_{14}	total idle time for machine
	ϕ_{15}	total idle time for all machines
	ϕ_{16}	current makespan
final makespan	ϕ_{17}	final makespan using SPT
	ϕ_{18}	final makespan using LPT
	ϕ_{19}	final makespan using LWR
	ϕ_{20}	final makespan using MWR
	ϕ_{RND}	final makespans using 100 random rollouts
	ϕ_{21}	mean for ϕ_{RND}
	ϕ_{22}	standard deviation for ϕ_{RND}
	ϕ_{23}	minimum value for ϕ_{RND}
	ϕ_{24}	maximum value for ϕ_{RND}

Trajectory strategies for Φ

ALICE

Helga

Introduction

Problem
space

Subspace of
instances

Feature space

Algorithm
space

Performance
space

Footprints in
instance space

Preference set

Preference
learning

Conclusions

Following the **policy**:

- ★ (Φ^{OPT}) expert π_* .
- ★ $(\Phi^{ES, \rho})$ the policy obtained by optimising with CMA-ES.
- ★ (Φ^{SPT}) shortest processing time (SPT).
- ★ (Φ^{LPT}) longest processing time (LPT).
- ★ (Φ^{LWR}) least work remaining (LWR).
- ★ (Φ^{MWR}) most work remaining (MWR).
- ★ (Φ^{RND}) random policy (RND).
- ★ (Φ^{ALL}) union of all of the above.

Trajectory strategies for Φ

ALICE

Helga

Introduction

Problem
space

Subspace of
instances

Feature space

Algorithm
space

Performance
space

Footprints in
instance space

Preference set

Preference
learning

Conclusions

Following the **policy**:

- ★ (Φ^{OPT}) expert π_* .
- ★ $(\Phi^{ES, \rho})$ the policy obtained by optimising with CMA-ES.
- ★ (Φ^{SPT}) shortest processing time (SPT).
- ★ (Φ^{LPT}) longest processing time (LPT).
- ★ (Φ^{LWR}) least work remaining (LWR).
- ★ (Φ^{MWR}) most work remaining (MWR).
- ★ (Φ^{RND}) random policy (RND).
- ★ (Φ^{ALL}) union of all of the above.

Trajectory strategies for Φ

ALICE

Helga

Introduction

Problem
space

Subspace of
instances

Feature space

Algorithm
space

Performance
space

Footprints in
instance space

Preference set

Preference
learning

Conclusions

Following the **policy**:

- ★ (Φ^{OPT}) expert π_* .
- ★ $(\Phi^{ES, \rho})$ the policy obtained by optimising with CMA-ES.
- ★ **(Φ^{SPT})** shortest processing time (SPT).
- ★ (Φ^{LPT}) longest processing time (LPT).
- ★ (Φ^{LWR}) least work remaining (LWR).
- ★ (Φ^{MWR}) most work remaining (MWR).
- ★ (Φ^{RND}) random policy (RND).
- ★ (Φ^{ALL}) union of all of the above.

Trajectory strategies for Φ

ALICE

Helga

Introduction

Problem
space

Subspace of
instances

Feature space

Algorithm
space

Performance
space

Footprints in
instance space

Preference set

Preference
learning

Conclusions

Following the **policy**:

- ★ (Φ^{OPT}) expert π_* .
- ★ $(\Phi^{ES.\rho})$ the policy obtained by optimising with CMA-ES.
- ★ (Φ^{SPT}) shortest processing time (SPT).
- ★ **(Φ^{LPT})** longest processing time (LPT).
- ★ (Φ^{LWR}) least work remaining (LWR).
- ★ (Φ^{MWR}) most work remaining (MWR).
- ★ (Φ^{RND}) random policy (RND).
- ★ (Φ^{ALL}) union of all of the above.

Trajectory strategies for Φ

ALICE

Helga

Introduction

Problem
space

Subspace of
instances

Feature space

Algorithm
space

Performance
space

Footprints in
instance space

Preference set

Preference
learning

Conclusions

Following the **policy**:

- ★ (Φ^{OPT}) expert π_* .
- ★ $(\Phi^{ES, \rho})$ the policy obtained by optimising with CMA-ES.
- ★ (Φ^{SPT}) shortest processing time (SPT).
- ★ (Φ^{LPT}) longest processing time (LPT).
- ★ (Φ^{LWR}) least work remaining (LWR).
- ★ (Φ^{MWR}) most work remaining (MWR).
- ★ (Φ^{RND}) random policy (RND).
- ★ (Φ^{ALL}) union of all of the above.

Trajectory strategies for Φ

ALICE

Helga

Introduction

Problem
space

Subspace of
instances

Feature space

Algorithm
space

Performance
space

Footprints in
instance space

Preference set

Preference
learning

Conclusions

Following the **policy**:

- ★ (Φ^{OPT}) expert π_* .
- ★ $(\Phi^{ES.\rho})$ the policy obtained by optimising with CMA-ES.
- ★ (Φ^{SPT}) shortest processing time (SPT).
- ★ (Φ^{LPT}) longest processing time (LPT).
- ★ (Φ^{LWR}) least work remaining (LWR).
- ★ (Φ^{MWR}) most work remaining (MWR).
- ★ (Φ^{RND}) random policy (RND).
- ★ (Φ^{ALL}) union of all of the above.

Trajectory strategies for Φ

ALICE

Helga

Introduction

Problem
space

Subspace of
instances

Feature space

Algorithm
space

Performance
space

Footprints in
instance space

Preference set

Preference
learning

Conclusions

Following the **policy**:

- ★ (Φ^{OPT}) expert π_* .
- ★ $(\Phi^{ES, \rho})$ the policy obtained by optimising with CMA-ES.
- ★ (Φ^{SPT}) shortest processing time (SPT).
- ★ (Φ^{LPT}) longest processing time (LPT).
- ★ (Φ^{LWR}) least work remaining (LWR).
- ★ (Φ^{MWR}) most work remaining (MWR).
- ★ **(Φ^{RND})** random policy (RND).
- ★ (Φ^{ALL}) union of all of the above.

Trajectory strategies for Φ

ALICE

Helga

Introduction

Problem
space

Subspace of
instances

Feature space

Algorithm
space

Performance
space

Footprints in
instance space

Preference set

Preference
learning

Conclusions

Following the policy:

- ★ (Φ^{OPT}) expert π_* .
- ★ $(\Phi^{ES.\rho})$ the policy obtained by optimising with CMA-ES.
- ★ (Φ^{SPT}) shortest processing time (SPT).
- ★ (Φ^{LPT}) longest processing time (LPT).
- ★ (Φ^{LWR}) least work remaining (LWR).
- ★ (Φ^{MWR}) most work remaining (MWR).
- ★ (Φ^{RND}) random policy (RND).
- ★ (Φ^{ALL}) union of all of the above.

Sampled size of $|\Phi|$ (6×5 , $N_{train} = 500$)

ALICE

Helga

Introduction

Problem space

Subspace of instances

Feature space

Algorithm space


Performance space

Footprints in instance space

Preference set

Preference learning

Conclusions



Framework for Algorithm Learning

ALICE

Helga

Introduction

Problem
space

Subspace of
instances

Feature space

Algorithm
space

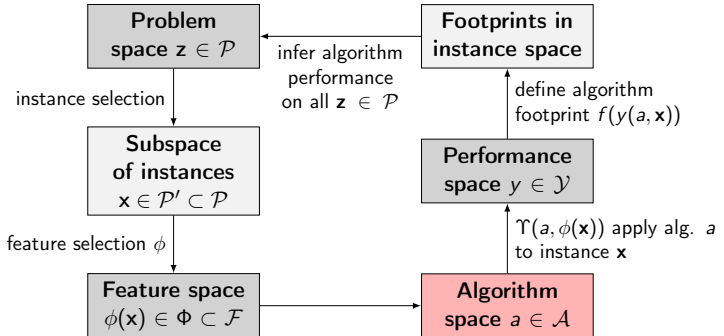
Performance
space

Footprints in
instance space

Preference set

Preference
learning

Conclusions





Various Methods for Solving JSP

ALICE

Helga

Introduction

Problem space

Subspace of instances

Feature space

Algorithm space

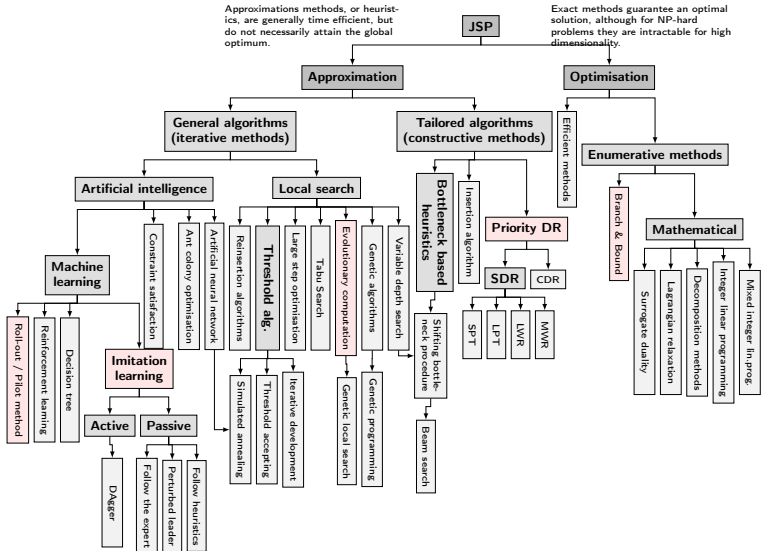
Performance space

Footprints in instance space

Preference set

Preference learning

Conclusions



Framework for Algorithm Learning

ALICE

Helga

Introduction

Problem
space

Subspace of
instances

Feature space

Algorithm
space

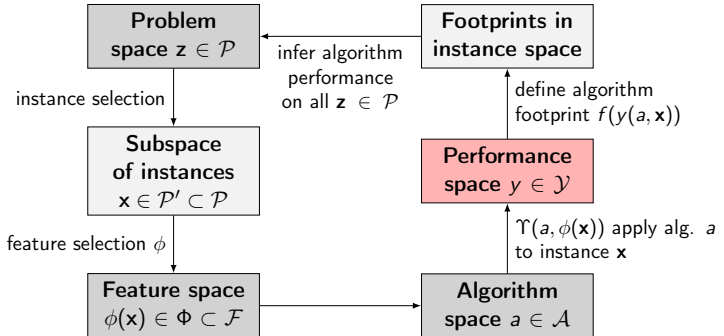
Performance
space

Footprints in
instance space

Preference set

Preference
learning

Conclusions





Performance measure

ALICE

Helga

Introduction

Problem
space

Subspace of
instances

Feature space

Algorithm
space

Performance
space

Footprints in
instance space

Preference set

Preference
learning

Conclusions

Performance of policy π compared with its optimal makespan, found using an expert policy, π_* , is the following loss function:

$$\rho = \frac{C_{\max}^{\pi} - C_{\max}^{\pi_*}}{C_{\max}^{\pi_*}} \cdot 100\%$$

The goal is to minimise this discrepancy between **predicted** value and **true** outcome.

Framework for Algorithm Learning

ALICE

Helga

Introduction

Problem
space

Subspace of
instances

Feature space

Algorithm
space

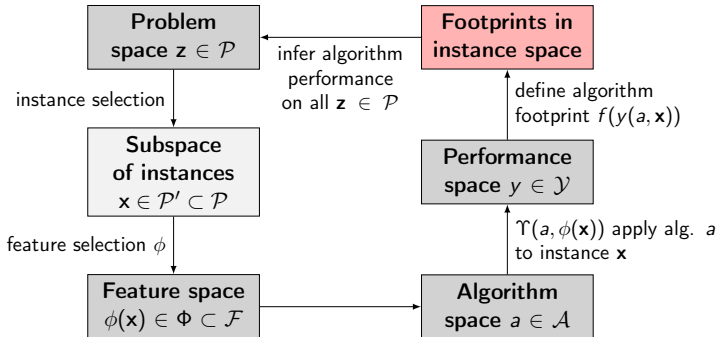
Performance
space

Footprints in
instance space

Preference set

Preference
learning

Conclusions



Deviation from optimality, ρ

ALICE

Helga

Introduction

Problem space

Subspace of instances

Feature space

Algorithm space



Performance space

Footprints in instance space

Preference set

Preference learning

Conclusions



Simple dispatching rule

- Shortest Processing Time
- Longest Processing Time
- Least Work Remaining
- Most Work Remaining
- Random dispatches

Data set



Deviation from optimality, ρ

ALICE

Helga

Introduction

Problem space

Subspace of instances

Feature space

Algorithm space



Performance space

Footprints in instance space

Preference set

Preference learning

Conclusions



Simple dispatching rule

- Shortest Processing Time
- Longest Processing Time
- Least Work Remaining
- Most Work Remaining
- Random dispatches

Data set

- train
- test

Making optimal decisions, ξ_{π}^{\star}

ALICE

Helga

Introduction

Problem
space

Subspace of
instances

Feature space

Algorithm
space


Performance
space

Footprints in
instance space

Preference set

Preference
learning

Conclusions



Making optimal decisions, ξ_π^*

ALICE

Helga

Introduction

Problem
space

Subspace of
instances

Feature space

Algorithm
space


Performance
space

Footprints in
instance space

Preference set

Preference
learning

Conclusions



Making optimal decisions, ξ_π

ALICE

Helga

Introduction

Problem
space

Subspace of
instances

Feature space

Algorithm
space


Performance
space

Footprints in
instance space

Preference set

Preference
learning

Conclusions



SDR — LPT — LWR — MWR — RND — SPT Accuracy — ξ^* -- ξ

Impact of suboptimal decision, $\{\varsigma_{\min}^*, \varsigma_{\max}^*\}$

ALICE

Helga

Introduction

Problem space

Subspace of instances

Feature space

Algorithm space


Performance space

Footprints in instance space

Preference set

Preference learning

Conclusions



Impact of suboptimal decision, $\{\zeta_{\min}^{\pi}, \zeta_{\max}^{\pi}\}$

ALICE

Helga

Introduction

Problem space

Subspace of instances

Feature space

Algorithm space


Performance space

Footprints in instance space

Preference set

Preference learning

Conclusions



Blended dispatching rule

ALICE

Helga

Introduction

Problem
space

Subspace of
instances

Feature space

Algorithm
space


Performance
space

Footprints in
instance space

Preference set

Preference
learning

Conclusions



Framework for Algorithm Learning

ALICE

Helga

Introduction

Problem
space

Subspace of
instances

Feature space

Algorithm
space

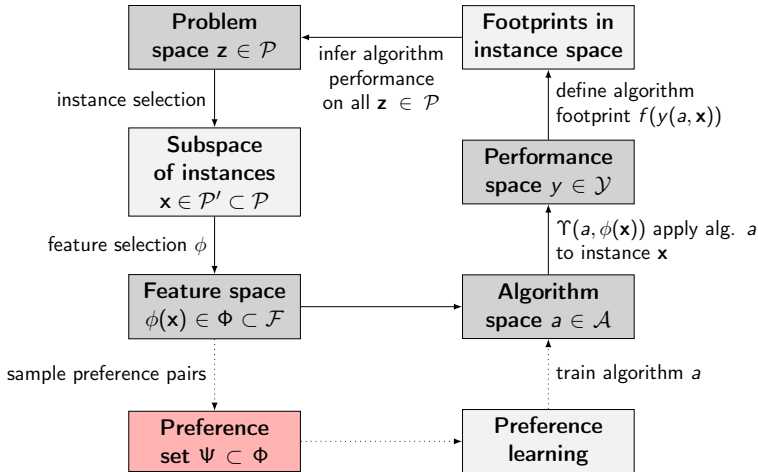
Performance
space

Footprints in
instance space

Preference set

Preference
learning

Conclusions



Generating training data

ALICE

Helga

Introduction

Problem
space

Subspace of
instances

Feature space

Algorithm
space

Performance
space

Footprints in
instance space

Preference set

Preference
learning

Conclusions

ALICE framework for creating dispatching rules:

- ★ **Linear classification** to identify good dispatches, from worse ones.
- ★ Generate feature set, $\Phi \subset \mathcal{F}$, both from
 - ★ optimal solutions, ϕ^o
 - ★ suboptimal solutions, ϕ^sby exploring various trajectories within the feature-space (where $\phi^o, \phi^s \in \mathcal{F}$).
- ★ Sample Φ to create training set Ψ with rank pairs:
 - ★ optimal decision, $(z^o, y_o) = (\phi^o - \phi^s, +1)$
 - ★ suboptimal decision, $(z^s, y_s) = (\phi^s - \phi^o, -1)$using different ranking schemes (where $z^o, z^s \in \Psi$)
- ★ Sample Ψ using stepwise bias for time independent policy.

Generating training data

ALICE

Helga

Introduction

Problem
space

Subspace of
instances

Feature space

Algorithm
space

Performance
space

Footprints in
instance space

Preference set

Preference
learning

Conclusions

ALICE framework for creating dispatching rules:

- ★ Linear classification to identify good dispatches, from worse ones.
- ★ **Generate** feature set, $\Phi \subset \mathcal{F}$, both from
 - ★ **optimal** solutions, ϕ^o
 - ★ **suboptimal** solutions, ϕ^sby exploring various **trajectories** within the feature-space (where $\phi^o, \phi^s \in \mathcal{F}$).
- ★ Sample Φ to create training set Ψ with rank pairs:
 - ★ optimal decision, $(z^o, y_o) = (\phi^o - \phi^s, +1)$
 - ★ suboptimal decision, $(z^s, y_s) = (\phi^s - \phi^o, -1)$using different ranking schemes (where $z^o, z^s \in \Psi$)
- ★ Sample Ψ using stepwise bias for time independent policy.

Generating training data

ALICE

Helga

Introduction

Problem
space

Subspace of
instances

Feature space

Algorithm
space

Performance
space

Footprints in
instance space

Preference set

Preference
learning

Conclusions

ALICE framework for creating dispatching rules:

- ★ Linear classification to identify good dispatches, from worse ones.
- ★ Generate feature set, $\Phi \subset \mathcal{F}$, both from
 - ★ optimal solutions, ϕ^o
 - ★ suboptimal solutions, ϕ^sby exploring various trajectories within the feature-space (where $\phi^o, \phi^s \in \mathcal{F}$).
- ★ Sample Φ to **create** training set Ψ with rank pairs:
 - ★ **optimal** decision, $(z^o, y_o) = (\phi^o - \phi^s, +1)$
 - ★ **suboptimal** decision, $(z^s, y_s) = (\phi^s - \phi^o, -1)$using different **ranking** schemes (where $z^o, z^s \in \Psi$)
- ★ Sample Ψ using stepwise bias for time independent policy.

Generating training data

ALICE

Helga

Introduction

Problem
space

Subspace of
instances

Feature space

Algorithm
space

Performance
space

Footprints in
instance space

Preference set

Preference
learning

Conclusions

ALICE framework for creating dispatching rules:

- ★ Linear classification to identify good dispatches, from worse ones.
- ★ Generate feature set, $\Phi \subset \mathcal{F}$, both from
 - ★ optimal solutions, ϕ^o
 - ★ suboptimal solutions, ϕ^sby exploring various trajectories within the feature-space (where $\phi^o, \phi^s \in \mathcal{F}$).
- ★ Sample Φ to create training set Ψ with rank pairs:
 - ★ optimal decision, $(z^o, y_o) = (\phi^o - \phi^s, +1)$
 - ★ suboptimal decision, $(z^s, y_s) = (\phi^s - \phi^o, -1)$using different ranking schemes (where $z^o, z^s \in \Psi$)
- ★ **Sample Ψ** using **stepwise bias** for time independent policy.

Ranking schemes for Ψ

ALICE

Helga

Introduction

Problem
space

Subspace of
instances

Feature space

Algorithm
space

Performance
space

Footprints in
instance space

Preference set

Preference
learning

Conclusions

Sampling rankings of available jobs where where

$r_1 > r_2 > \dots > r_{n'} \ (n' \leq n)$ with respect to:

Ψ_b all opt rankings r_1 vs. all possible subopt rankings r_i ,
 $i \in \{2, \dots, n'\}$

Ψ_f full subsequent rankings, i.e., all combinations of r_i and
 r_{i+1} for all $i \in \{1, \dots, n'\}$.

Ψ_p partial subsequent rankings, similar to Ψ_f except if there
are more than one operation with the same rank, only one
is needed to be compared to subsequent rank ($\Psi_p \subset \Psi_f$).

Ψ_a union of all of the above.

Ranking schemes for Ψ

ALICE

Helga

Introduction

Problem
space

Subspace of
instances

Feature space

Algorithm
space

Performance
space

Footprints in
instance space

Preference set

Preference
learning

Conclusions

Sampling rankings of available jobs where where

$r_1 > r_2 > \dots > r_{n'} \ (n' \leq n)$ with respect to:

Ψ_b **all** opt rankings r_1 vs. all possible subopt rankings r_i ,
 $i \in \{2, \dots, n'\}$

Ψ_f full subsequent rankings, i.e., all combinations of r_i and
 r_{i+1} for all $i \in \{1, \dots, n'\}$.

Ψ_p partial subsequent rankings, similar to Ψ_f except if there
are more than one operation with the same rank, only one
is needed to be compared to subsequent rank ($\Psi_p \subset \Psi_f$).

Ψ_a union of all of the above.

Ranking schemes for Ψ

ALICE

Helga

Introduction

Problem
space

Subspace of
instances

Feature space

Algorithm
space

Performance
space

Footprints in
instance space

Preference set

Preference
learning

Conclusions

Sampling rankings of available jobs where where

$r_1 > r_2 > \dots > r_{n'} \ (n' \leq n)$ with respect to:

Ψ_b all opt rankings r_1 vs. all possible subopt rankings r_i ,
 $i \in \{2, \dots, n'\}$

Ψ_f **full subsequent** rankings, i.e., all combinations of r_i and
 r_{i+1} for all $i \in \{1, \dots, n'\}$.

Ψ_p partial subsequent rankings, similar to Ψ_f except if there
are more than one operation with the same rank, only one
is needed to be compared to subsequent rank ($\Psi_p \subset \Psi_f$).

Ψ_a union of all of the above.

Ranking schemes for Ψ

ALICE

Helga

Introduction

Problem
space

Subspace of
instances

Feature space

Algorithm
space

Performance
space

Footprints in
instance space

Preference set

Preference
learning

Conclusions

Sampling rankings of available jobs where where

$r_1 > r_2 > \dots > r_{n'} \ (n' \leq n)$ with respect to:

Ψ_b all opt rankings r_1 vs. all possible subopt rankings r_i ,
 $i \in \{2, \dots, n'\}$

Ψ_f full subsequent rankings, i.e., all combinations of r_i and
 r_{i+1} for all $i \in \{1, \dots, n'\}$.

Ψ_p **partial subsequent** rankings, similar to Ψ_f except if there
are more than one operation with the same rank, only one
is needed to be compared to subsequent rank ($\Psi_p \subset \Psi_f$).

Ψ_a union of all of the above.

Ranking schemes for Ψ

ALICE

Helga

Introduction

Problem
space

Subspace of
instances

Feature space

Algorithm
space

Performance
space

Footprints in
instance space

Preference set

Preference
learning

Conclusions

Sampling rankings of available jobs where where

$r_1 > r_2 > \dots > r_{n'} \ (n' \leq n)$ with respect to:

Ψ_b all opt rankings r_1 vs. all possible subopt rankings r_i ,
 $i \in \{2, \dots, n'\}$

Ψ_f full subsequent rankings, i.e., all combinations of r_i and
 r_{i+1} for all $i \in \{1, \dots, n'\}$.

Ψ_p partial subsequent rankings, similar to Ψ_f except if there
are more than one operation with the same rank, only one
is needed to be compared to subsequent rank ($\Psi_p \subset \Psi_f$).

Ψ_a union of **all** of the above.

Sampled size of $|\Psi|$ (6×5 , $N_{train} = 500$)

ALICE

Helga

Introduction

Problem
space

Subspace of
instances

Feature space

Algorithm
space


Performance
space

Footprints in
instance space

Preference set

Preference
learning

Conclusions



Stepwise bias for sampling ψ

ALICE

Helga

Introduction

Problem
space

Subspace of
instances

Feature space

Algorithm
space

Performance
space

Footprints in
instance space

Preference set

Preference
learning

Conclusions

Sampling **stepwise bias** for preference pairs:

- ★ (equal) equal probability.
- ★ (opt) inverse optimality for random dispatches $- 1 - \zeta_{\text{RND}}^*$.
- ★ (bcs) best case scenario for mean $\rho - \zeta_{\text{min}}^*$.
- ★ (wcs) worst case scenario for mean $\rho - \zeta_{\text{max}}^*$.
- ★ (featsize) inversely proportional to $|\Phi^{\text{OPT}}|$
- ★ (prefsize) inversely proportional to $|\Psi_{\rho}^{\text{OPT}}|$
- ★ (dbl1st) twice as much weight on the first half of the dispatches.
- ★ (dbl2nd) twice as much weight on the second half of the dispatches.

Stepwise bias for sampling ψ

ALICE

Helga

Introduction

Problem
space

Subspace of
instances

Feature space

Algorithm
space

Performance
space

Footprints in
instance space

Preference set

Preference
learning

Conclusions

Sampling **stepwise bias** for preference pairs:

- ★ (equal) equal probability.
- ★ (opt) inverse optimality for random dispatches – $1 - \zeta_{\text{RND}}^*$.
- ★ (bcs) best case scenario for mean $\rho - \zeta_{\text{min}}^*$.
- ★ (wcs) worst case scenario for mean $\rho - \zeta_{\text{max}}^*$.
- ★ (featsize) inversely proportional to $|\Phi^{\text{OPT}}|$
- ★ (prefsize) inversely proportional to $|\Psi_{\rho}^{\text{OPT}}|$
- ★ (dbl1st) twice as much weight on the first half of the dispatches.
- ★ (dbl2nd) twice as much weight on the second half of the dispatches.

Stepwise bias for sampling ψ

ALICE

Helga

Introduction

Problem
space

Subspace of
instances

Feature space

Algorithm
space

Performance
space

Footprints in
instance space

Preference set

Preference
learning

Conclusions

Sampling **stepwise bias** for preference pairs:

- ★ (equal) equal probability.
- ★ **(opt)** inverse optimality for random dispatches $-1 - \xi_{\text{RND}}^*$.
- ★ (bcs) best case scenario for mean $\rho - \zeta_{\text{min}}^*$.
- ★ (wcs) worst case scenario for mean $\rho - \zeta_{\text{max}}^*$.
- ★ (featsize) inversely proportional to $|\Phi^{\text{OPT}}|$
- ★ (prefsize) inversely proportional to $|\Psi_{\rho}^{\text{OPT}}|$
- ★ (dbl1st) twice as much weight on the first half of the dispatches.
- ★ (dbl2nd) twice as much weight on the second half of the dispatches.



Stepwise bias for sampling ψ

ALICE

Helga

Introduction

Problem
space

Subspace of
instances

Feature space

Algorithm
space

Performance
space

Footprints in
instance space

Preference set

Preference
learning

Conclusions

Sampling **stepwise bias** for preference pairs:

- ★ (equal) equal probability.
- ★ (opt) inverse optimality for random dispatches $-1 - \zeta_{\text{RND}}^*$.
- ★ **(bcs)** best case scenario for mean $\rho - \zeta_{\text{min}}^*$.
- ★ (wcs) worst case scenario for mean $\rho - \zeta_{\text{max}}^*$.
- ★ (featsize) inversely proportional to $|\Phi^{\text{OPT}}|$
- ★ (prefsize) inversely proportional to $|\Psi_{\rho}^{\text{OPT}}|$
- ★ (dbl1st) twice as much weight on the first half of the dispatches.
- ★ (dbl2nd) twice as much weight on the second half of the dispatches.

Stepwise bias for sampling ψ

ALICE

Helga

Introduction

Problem
space

Subspace of
instances

Feature space

Algorithm
space

Performance
space

Footprints in
instance space

Preference set

Preference
learning

Conclusions

Sampling **stepwise bias** for preference pairs:

- ★ (equal) equal probability.
- ★ (opt) inverse optimality for random dispatches $-1 - \xi_{\text{RND}}^*$.
- ★ (bcs) best case scenario for mean $\rho - \zeta_{\text{min}}^*$.
- ★ **(wcs)** worst case scenario for mean $\rho - \zeta_{\text{max}}^*$.
- ★ (featsize) inversely proportional to $|\Phi^{\text{OPT}}|$
- ★ (prefsize) inversely proportional to $|\Psi_{\rho}^{\text{OPT}}|$
- ★ (dbl1st) twice as much weight on the first half of the dispatches.
- ★ (dbl2nd) twice as much weight on the second half of the dispatches.



Stepwise bias for sampling ψ

ALICE

Helga

Introduction

Problem
space

Subspace of
instances

Feature space

Algorithm
space

Performance
space

Footprints in
instance space

Preference set

Preference
learning

Conclusions

Sampling **stepwise bias** for preference pairs:

- ★ (equal) equal probability.
- ★ (opt) inverse optimality for random dispatches $-1 - \xi_{\text{RND}}^*$.
- ★ (bcs) best case scenario for mean $\rho - \zeta_{\text{min}}^*$.
- ★ (wcs) worst case scenario for mean $\rho - \zeta_{\text{max}}^*$.
- ★ (**featsize**) inversely proportional to $|\Phi^{\text{OPT}}|$
- ★ (prefsize) inversely proportional to $|\Psi_{\rho}^{\text{OPT}}|$
- ★ (dbl1st) twice as much weight on the first half of the dispatches.
- ★ (dbl2nd) twice as much weight on the second half of the dispatches.

Stepwise bias for sampling ψ

ALICE

Helga

Introduction

Problem
space

Subspace of
instances

Feature space

Algorithm
space

Performance
space

Footprints in
instance space

Preference set

Preference
learning

Conclusions

Sampling **stepwise bias** for preference pairs:

- ★ (equal) equal probability.
- ★ (opt) inverse optimality for random dispatches $-1 - \xi_{\text{RND}}^*$.
- ★ (bcs) best case scenario for mean $\rho - \zeta_{\text{min}}^*$.
- ★ (wcs) worst case scenario for mean $\rho - \zeta_{\text{max}}^*$.
- ★ (featsize) inversely proportional to $|\Phi^{\text{OPT}}|$
- ★ **(prefsize)** inversely proportional to $|\Psi_{\rho}^{\text{OPT}}|$
- ★ (dbl1st) twice as much weight on the first half of the dispatches.
- ★ (dbl2nd) twice as much weight on the second half of the dispatches.

Stepwise bias for sampling ψ

ALICE

Helga

Introduction

Problem
space

Subspace of
instances

Feature space

Algorithm
space

Performance
space

Footprints in
instance space

Preference set

Preference
learning

Conclusions

Sampling **stepwise bias** for preference pairs:

- ★ (equal) equal probability.
- ★ (opt) inverse optimality for random dispatches $-1 - \xi_{\text{RND}}^*$.
- ★ (bcs) best case scenario for mean $\rho - \zeta_{\text{min}}^*$.
- ★ (wcs) worst case scenario for mean $\rho - \zeta_{\text{max}}^*$.
- ★ (featsize) inversely proportional to $|\Phi^{\text{OPT}}|$
- ★ (prefsize) inversely proportional to $|\Psi_{\rho}^{\text{OPT}}|$
- ★ **(dbl1st)** twice as much weight on the first half of the dispatches.
- ★ (dbl2nd) twice as much weight on the second half of the dispatches.



Stepwise bias for sampling ψ

ALICE

Helga

Introduction

Problem
space

Subspace of
instances

Feature space

Algorithm
space

Performance
space

Footprints in
instance space

Preference set

Preference
learning

Conclusions

Sampling **stepwise bias** for preference pairs:

- ★ (equal) equal probability.
- ★ (opt) inverse optimality for random dispatches $-1 - \xi_{\text{RND}}^*$.
- ★ (bcs) best case scenario for mean $\rho - \zeta_{\text{min}}^*$.
- ★ (wcs) worst case scenario for mean $\rho - \zeta_{\text{max}}^*$.
- ★ (featsize) inversely proportional to $|\Phi^{\text{OPT}}|$
- ★ (prefsize) inversely proportional to $|\Psi_p^{\text{OPT}}|$
- ★ (dbl1st) twice as much weight on the first half of the dispatches.
- ★ (**dbl2nd**) twice as much weight on the second half of the dispatches.

Stepwise bias strategies

ALICE

Helga

Introduction

Problem
space

Subspace of
instances

Feature space

Algorithm
space


Performance
space

Footprints in
instance space

Preference set

Preference
learning

Conclusions



Framework for Algorithm Learning

ALICE

Helga

Introduction

Problem space

Subspace of instances

Feature space

Algorithm space

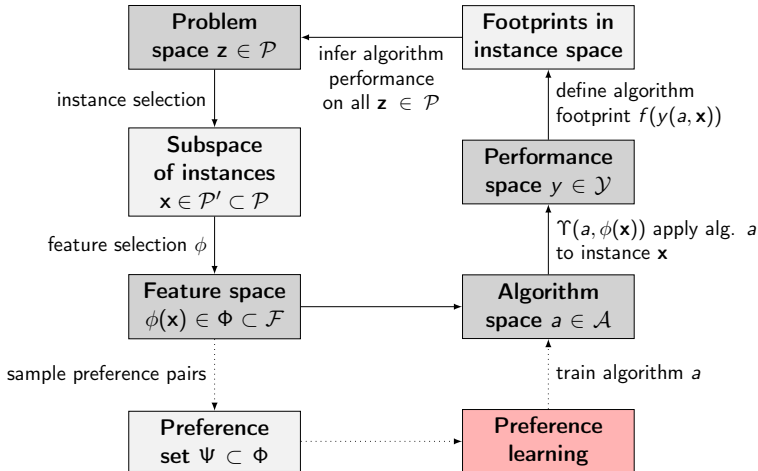
Performance space

Footprints in instance space

Preference set

Preference learning

Conclusions





Ordinal Regression

ALICE

Helga

Introduction

Problem
space

Subspace of
instances

Feature space

Algorithm
space

Performance
space

Footprints in
instance space

Preference set

Preference
learning

Conclusions

Preference learning:

- ★ Mapping of points to ranks: $\{h(\cdot) : \Phi \mapsto Y\}$ where

$$\phi_o \succ \phi_s \iff h(\phi_o) > h(\phi_s)$$

- ★ The preference is defined by a linear function:

$$h(\phi) = \sum_{i=1}^d w_i \phi_i$$

optimised w.r.t. w based on training data Ψ

- ★ Note: Limitations in approximation function to capture the complex dynamics incorporated in optimal trajectories.

Ordinal Regression

ALICE

Helga

Introduction

Problem
space

Subspace of
instances

Feature space

Algorithm
space

Performance
space

Footprints in
instance space

Preference set

Preference
learning

Conclusions

Preference **learning**:

- ★ Mapping of points to ranks: $\{h(\cdot) : \Phi \mapsto Y\}$ where

$$\phi_o \succ \phi_s \iff h(\phi_o) > h(\phi_s)$$

- ★ The preference is defined by a **linear** function:

$$h(\phi) = \sum_{i=1}^d w_i \phi$$

optimised w.r.t. \mathbf{w} based on training data Ψ

- ★ Note: Limitations in approximation function to capture the complex dynamics incorporated in optimal trajectories.

Ordinal Regression

ALICE

Helga

Introduction

Problem
space

Subspace of
instances

Feature space

Algorithm
space

Performance
space

Footprints in
instance space

Preference set

Preference
learning

Conclusions

Preference learning:

- ★ Mapping of points to ranks: $\{h(\cdot) : \Phi \mapsto Y\}$ where

$$\phi_o \succ \phi_s \iff h(\phi_o) > h(\phi_s)$$

- ★ The preference is defined by a **linear** function:

$$h(\phi) = \sum_{i=1}^d w_i \phi$$

optimised w.r.t. \mathbf{w} based on training data Ψ

- ★ Note: **Limitations** in **approximation** function to capture the complex dynamics incorporated in optimal trajectories.

Learning methods

ALICE

Helga

Introduction

Problem
space

Subspace of
instances

Feature space

Algorithm
space

Performance
space

Footprints in
instance space

Preference set

Preference
learning

Conclusions

Passive imitation learning (single pass):

- ★ Prediction with expert advice, π_\star – Gurobi
- ★ Follow the perturbed leader (OPT_ϵ).

Active imitation learning (iterative):

- ★ Dataset Aggregation (DAgger)

$$\pi_i = \beta_i \pi_\star + (1 - \beta_i) \hat{\pi}_{i-1}$$

where $\hat{\pi}_{i-1}$ is the previous learned model, and $\hat{\pi}_i$ learns on

$$\phi^{\text{DA}i} = \bigcup_{i'=0}^i \phi^{\text{IL}i'}$$

aggregated dataset of all previous iterations.

Learning methods

ALICE

Helga

Introduction

Problem
space

Subspace of
instances

Feature space

Algorithm
space

Performance
space

Footprints in
instance space

Preference set

Preference
learning

Conclusions

Passive imitation learning (single pass):

- ★ Prediction with expert advice, π_\star – Gurobi
- ★ Follow the perturbed leader (**OPT** ϵ).

Active imitation learning (iterative):

- ★ Dataset Aggregation (DAgger)

$$\pi_i = \beta_i \pi_\star + (1 - \beta_i) \hat{\pi}_{i-1}$$

where $\hat{\pi}_{i-1}$ is the previous learned model, and $\hat{\pi}_i$ learns on

$$\phi^{\text{DA}i} = \bigcup_{i'=0}^i \phi^{\text{IL}i'}$$

aggregated dataset of all previous iterations.

Learning methods

ALICE

Helga

Introduction

Problem
space

Subspace of
instances

Feature space

Algorithm
space

Performance
space

Footprints in
instance space

Preference set

Preference
learning

Conclusions

Passive imitation learning (single pass):

- ★ Prediction with expert advice, π_{\star} – Gurobi
- ★ Follow the perturbed leader (OPT_{ϵ}).

Active imitation learning (iterative):

- ★ Dataset Aggregation (DAgger)

$$\pi_i = \beta_i \pi_{\star} + (1 - \beta_i) \hat{\pi}_{i-1}$$

where $\hat{\pi}_{i-1}$ is the previous learned model, and $\hat{\pi}_i$ learns on

$$\phi^{\text{DA}i} = \bigcup_{i'=0}^i \phi^{\text{IL}i'}$$

aggregated dataset of all previous iterations.

Learning methods

ALICE

Helga

Introduction

Problem
space

Subspace of
instances

Feature space

Algorithm
space

Performance
space

Footprints in
instance space

Preference set

Preference
learning

Conclusions

Passive imitation learning (single pass):

- ★ Prediction with expert advice, π_{\star} – Gurobi
- ★ Follow the perturbed leader (OPT_{ϵ}).

Active imitation learning (iterative):

- ★ Dataset Aggregation (DAgger)

$$\pi_i = \beta_i \pi_{\star} + (1 - \beta_i) \hat{\pi}_{i-1}$$

where $\hat{\pi}_{i-1}$ is the previous learned model, and $\hat{\pi}_i$ learns on

$$\phi^{\text{DA}i} = \bigcup_{i'=0}^i \phi^{\text{IL}i'}$$

aggregated dataset of all previous iterations.

ALICE

Helga

Introduction

Problem
space

Subspace of
instances

Feature space

Algorithm
space

Performance
space

Footprints in
instance space

Preference set

Preference
learning

Conclusions

Require: $T \geq 1$

```
1: procedure DAgger( $\pi_*, \Phi^{\text{IL}0}, T$ )
2:    $\hat{\pi}_0 \leftarrow \text{Train}(\Phi^{\text{IL}0})$   $\triangleright$  initial model, iff  $\Phi^{\text{IL}0} = \Phi^{\text{OPT}}$ 
3:   for  $i \leftarrow 1$  to  $T$  do  $\triangleright$  at each imitation learning iteration
4:     Let  $\pi_i = \beta_i \pi_* + (1 - \beta_i) \hat{\pi}_{i-1}$ 
5:     Sample a  $K$ -solution using  $\pi_i$   $\triangleright \text{IL}(i, \hat{\pi}_{i-1}, \pi_*)$ 
6:      $\Phi^{\text{IL}i} = \{(s, \pi_*(s))\}$   $\triangleright$  visited by  $\pi_i$  and actions by  $\pi_*$ 
7:      $\Phi^{\text{DA}i} \leftarrow \Phi^{\text{DA}i-1} \cup \Phi^{\text{IL}i}$   $\triangleright$  aggregate datasets
8:      $\hat{\pi}_{i+1} \leftarrow \text{Train}(\Phi^{\text{DA}i})$   $\triangleright$  preference model
9:   end for
10:  return best  $\hat{\pi}_i$  on validation  $\triangleright$  best preference model
11: end procedure
```

Require: $T \geq 1$

- 1: **procedure** DAgger($\pi_*, \Phi^{\text{IL0}}, T$)
- 2: $\hat{\pi}_0 \leftarrow \text{Train}(\Phi^{\text{IL0}})$ \triangleright initial model, iff $\Phi^{\text{IL0}} = \Phi^{\text{OPT}}$
- 3: **for** $i \leftarrow 1$ **to** T **do** \triangleright at each imitation learning iteration
- 4: Let $\pi_i = \beta_i \pi_* + (1 - \beta_i) \hat{\pi}_{i-1}$
- 5: Sample a K -solution using π_i $\triangleright \text{IL}(i, \hat{\pi}_{i-1}, \pi_*)$
- 6: $\Phi^{\text{IL}i} = \{(s, \pi_*(s))\}$ \triangleright visited by π_i and actions by π_*
- 7: $\Phi^{\text{DA}i} \leftarrow \Phi^{\text{DA}i-1} \cup \Phi^{\text{IL}i}$ \triangleright aggregate datasets
- 8: $\hat{\pi}_{i+1} \leftarrow \text{Train}(\Phi^{\text{DA}i})$ \triangleright preference model
- 9: **end for**
- 10: **return** best $\hat{\pi}_i$ on validation \triangleright best preference model
- 11: **end procedure**



DAgger

ALICE

Helga

Introduction

Problem
space

Subspace of
instances

Feature space

Algorithm
space

Performance
space

Footprints in
instance space

Preference set

Preference
learning

Conclusions

Require: $T \geq 1$

- 1: **procedure** DAgger($\pi_*, \Phi^{\text{ILO}}, T$)
- 2: $\hat{\pi}_0 \leftarrow \text{Train}(\Phi^{\text{ILO}})$ \triangleright initial model, iff $\Phi^{\text{ILO}} = \Phi^{\text{OPT}}$
- 3: **for** $i \leftarrow 1$ **to** T **do** \triangleright at each imitation learning iteration
- 4: Let $\pi_i = \beta_i \pi_* + (1 - \beta_i) \hat{\pi}_{i-1}$
- 5: Sample a K -solution using π_i $\triangleright \text{IL}(i, \hat{\pi}_{i-1}, \pi_*)$
- 6: $\Phi^{\text{IL}i} = \{(s, \pi_*(s))\}$ \triangleright visited by π_i and actions by π_*
- 7: $\Phi^{\text{DA}i} \leftarrow \Phi^{\text{DA}i-1} \cup \Phi^{\text{IL}i}$ \triangleright aggregate datasets
- 8: $\hat{\pi}_{i+1} \leftarrow \text{Train}(\Phi^{\text{DA}i})$ \triangleright preference model
- 9: **end for**
- 10: **return** best $\hat{\pi}_i$ on validation \triangleright best preference model
- 11: **end procedure**

Require: $T \geq 1$

- 1: **procedure** DAgger($\pi_*, \Phi^{\text{ILO}}, T$)
- 2: $\hat{\pi}_0 \leftarrow \text{Train}(\Phi^{\text{ILO}})$ \triangleright initial model, iff $\Phi^{\text{ILO}} = \Phi^{\text{OPT}}$
- 3: **for** $i \leftarrow 1$ **to** T **do** \triangleright at each imitation learning iteration
- 4: Let $\pi_i = \beta_i \pi_* + (1 - \beta_i) \hat{\pi}_{i-1}$
- 5: Sample a K -solution using π_i $\triangleright \text{IL}(i, \hat{\pi}_{i-1}, \pi_*)$
- 6: $\Phi^{\text{IL}i} = \{(s, \pi_*(s))\}$ \triangleright visited by π_i and actions by π_*
- 7: $\Phi^{\text{DA}i} \leftarrow \Phi^{\text{DA}i-1} \cup \Phi^{\text{IL}i}$ \triangleright aggregate datasets
- 8: $\hat{\pi}_{i+1} \leftarrow \text{Train}(\Phi^{\text{DA}i})$ \triangleright preference model
- 9: **end for**
- 10: **return** best $\hat{\pi}_i$ on validation \triangleright best preference model
- 11: **end procedure**

Require: $T \geq 1$

```

1: procedure DAgger( $\pi_*$ ,  $\Phi^{\text{ILO}}$ ,  $T$ )
2:    $\hat{\pi}_0 \leftarrow \text{Train}(\Phi^{\text{ILO}})$   $\triangleright$  initial model, iff  $\Phi^{\text{ILO}} = \Phi^{\text{OPT}}$ 
3:   for  $i \leftarrow 1$  to  $T$  do  $\triangleright$  at each imitation learning iteration
4:     Let  $\pi_i = \beta_i \pi_* + (1 - \beta_i) \hat{\pi}_{i-1}$ 
5:     Sample a  $K$ -solution using  $\pi_i$   $\triangleright \text{IL}(i, \hat{\pi}_{i-1}, \pi_*)$ 
6:      $\Phi^{\text{LI}} = \{(s, \pi_*(s))\}$   $\triangleright$  visited by  $\pi_i$  and actions by  $\pi_*$ 
7:      $\Phi^{\text{DA}i} \leftarrow \Phi^{\text{DA}i-1} \cup \Phi^{\text{LI}}$   $\triangleright$  aggregate datasets
8:      $\hat{\pi}_{i+1} \leftarrow \text{Train}(\Phi^{\text{DA}i})$   $\triangleright$  preference model
9:   end for
10:  return best  $\hat{\pi}_i$  on validation  $\triangleright$  best preference model
11: end procedure
  
```


Require: $T \geq 1$

- 1: **procedure** DAgger($\pi_*, \Phi^{\text{ILO}}, T$)
- 2: $\hat{\pi}_0 \leftarrow \text{Train}(\Phi^{\text{ILO}})$ \triangleright initial model, iff $\Phi^{\text{ILO}} = \Phi^{\text{OPT}}$
- 3: **for** $i \leftarrow 1$ **to** T **do** \triangleright at each imitation learning iteration
- 4: Let $\pi_i = \beta_i \pi_* + (1 - \beta_i) \hat{\pi}_{i-1}$
- 5: Sample a K -solution using π_i $\triangleright \text{IL}(i, \hat{\pi}_{i-1}, \pi_*)$
- 6: $\Phi^{\text{IL}i} = \{(s, \pi_*(s))\}$ \triangleright visited by π_i and actions by π_*
- 7: $\Phi^{\text{DA}i} \leftarrow \Phi^{\text{DA}i-1} \cup \Phi^{\text{IL}i}$ \triangleright aggregate datasets
- 8: $\hat{\pi}_{i+1} \leftarrow \text{Train}(\Phi^{\text{DA}i})$ \triangleright preference model
- 9: **end for**
- 10: **return** best $\hat{\pi}_i$ on validation \triangleright best preference model
- 11: **end procedure**

Require: $T \geq 1$

```

1: procedure DAgger( $\pi_*$ ,  $\Phi^{\text{ILO}}, T$ )
2:    $\hat{\pi}_0 \leftarrow \text{Train}(\Phi^{\text{ILO}})$   $\triangleright$  initial model, iff  $\Phi^{\text{ILO}} = \Phi^{\text{OPT}}$ 
3:   for  $i \leftarrow 1$  to  $T$  do  $\triangleright$  at each imitation learning iteration
4:     Let  $\pi_i = \beta_i \pi_* + (1 - \beta_i) \hat{\pi}_{i-1}$ 
5:     Sample a  $K$ -solution using  $\pi_i$   $\triangleright \text{IL}(i, \hat{\pi}_{i-1}, \pi_*)$ 
6:      $\Phi^{\text{IL}i} = \{(s, \pi_*(s))\}$   $\triangleright$  visited by  $\pi_i$  and actions by  $\pi_*$ 
7:      $\Phi^{\text{DA}i} \leftarrow \Phi^{\text{DA}i-1} \cup \Phi^{\text{IL}i}$   $\triangleright$  aggregate datasets
8:      $\hat{\pi}_{i+1} \leftarrow \text{Train}(\Phi^{\text{DA}i})$   $\triangleright$  preference model
9:   end for
10:  return best  $\hat{\pi}_i$  on validation  $\triangleright$  best preference model
11: end procedure

```

Require: $T \geq 1$

- 1: **procedure** DAgger(π_* , Φ^{ILO}, T)
- 2: $\hat{\pi}_0 \leftarrow \text{Train}(\Phi^{\text{ILO}})$ \triangleright initial model, iff $\Phi^{\text{ILO}} = \Phi^{\text{OPT}}$
- 3: **for** $i \leftarrow 1$ **to** T **do** \triangleright at each imitation learning iteration
- 4: Let $\pi_i = \beta_i \pi_* + (1 - \beta_i) \hat{\pi}_{i-1}$
- 5: Sample a K -solution using π_i $\triangleright \text{IL}(i, \hat{\pi}_{i-1}, \pi_*)$
- 6: $\Phi^{\text{IL}i} = \{(s, \pi_*(s))\}$ \triangleright visited by π_i and actions by π_*
- 7: $\Phi^{\text{DA}i} \leftarrow \Phi^{\text{DA}i-1} \cup \Phi^{\text{IL}i}$ \triangleright aggregate datasets
- 8: $\hat{\pi}_{i+1} \leftarrow \text{Train}(\Phi^{\text{DA}i})$ \triangleright preference model
- 9: **end for**
- 10: **return** best $\hat{\pi}_i$ on validation \triangleright best preference model
- 11: **end procedure**

Deviation from optimality, ρ

ALICE

Helga

Introduction

Problem space

Subspace of instances

Feature space

Algorithm space


Performance space

Footprints in instance space

Preference set

Preference learning

Conclusions



Framework for Algorithm Learning

ALICE

Helga

Introduction

Problem space

Subspace of instances

Feature space

Algorithm space

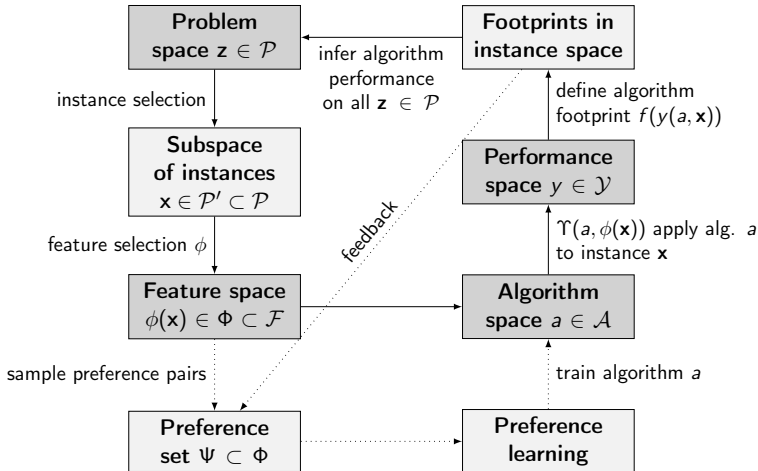
Performance space

Footprints in instance space

Preference set

Preference learning

Conclusions





Using Analysis & Learning Iterative Consecutive Executions framework I

ALICE

Helga

Introduction

Problem space

Subspace of instances

Feature space

Algorithm space

Performance space

Footprints in instance space

Preference set

Preference learning

Conclusions

The thesis introduced a framework for learning (**linear**) composite priority dispatching rule – using **job-shop** as a case-study – with the following guidelines:

- ★ For a given problem domain, use a suitable problem generator to train and test on.
- ★ Define features to grasp the essence of visited k -solutions
- ★ Success is highly dependent on the preference pairs introduced to the system:

- ★ Ψ_p reduces the preference set without loss of performance.
- ★ Stepwise bias is needed to balance time dependent Ψ_p in order to create time independent models.

It is non intuitive how to go about collecting training data.



Using Analysis & Learning Iterative Consecutive Executions framework I

ALICE

Helga

Introduction

Problem space

Subspace of instances

Feature space

Algorithm space

Performance space

Footprints in instance space

Preference set

Preference learning

Conclusions

The thesis introduced a framework for learning (linear) composite priority dispatching rule – using job-shop as a case-study – with the following guidelines:

- ★ For a given problem domain, use a suitable problem **generator** to **train** and **test** on.
- ★ Define features to grasp the essence of visited k -solutions
- ★ Success is highly dependent on the preference pairs introduced to the system:
 - ★ Ψ_p reduces the preference set without loss of performance.
 - ★ Stepwise bias is needed to balance time dependent Ψ_p in order to create time independent models.

It is non intuitive how to go about collecting training data.



Using Analysis & Learning Iterative Consecutive Executions framework I

ALICE

Helga

Introduction

Problem space

Subspace of instances

Feature space

Algorithm space

Performance space

Footprints in instance space

Preference set

Preference learning

Conclusions

The thesis introduced a framework for learning (linear) composite priority dispatching rule – using job-shop as a case-study – with the following guidelines:

- ★ For a given problem domain, use a suitable problem generator to train and test on.
- ★ Define **features** to grasp the essence of visited **k-solutions**
- ★ Success is highly dependent on the preference pairs introduced to the system:

- ★ Ψ_p reduces the preference set without loss of performance.
- ★ Stepwise bias is needed to balance time dependent Ψ_p in order to create time independent models.

It is non intuitive how to go about collecting training data.



Using Analysis & Learning Iterative Consecutive Executions framework I

ALICE

Helga

Introduction

Problem space

Subspace of instances

Feature space

Algorithm space

Performance space

Footprints in instance space

Preference set

Preference learning

Conclusions

The thesis introduced a framework for learning (linear) composite priority dispatching rule – using job-shop as a case-study – with the following guidelines:

- ★ For a given problem domain, use a suitable problem generator to train and test on.
- ★ Define features to grasp the essence of visited k -solutions
- ★ **Success** is highly dependent on the preference pairs introduced to the system:

- ★ Ψ_p reduces the preference set without loss of performance.

- ★ Stepwise bias is needed to balance time dependent Ψ_p in order to create time independent models.

It is non intuitive how to go about collecting training data.



Using Analysis & Learning Iterative Consecutive Executions framework I

ALICE

Helga

Introduction

Problem space

Subspace of instances

Feature space

Algorithm space

Performance space

Footprints in instance space

Preference set

Preference learning

Conclusions

The thesis introduced a framework for learning (linear) composite priority dispatching rule – using job-shop as a case-study – with the following guidelines:

- ★ For a given problem domain, use a suitable problem generator to train and test on.
 - ★ Define features to grasp the essence of visited k -solutions
 - ★ **Success** is highly dependent on the preference pairs introduced to the system:
 - ★ Ψ_p **reduces** the preference set without loss of performance.
 - ★ Stepwise bias is needed to balance time dependent Ψ_p in order to create time independent models.
- It is non intuitive how to go about collecting training data.



Using Analysis & Learning Iterative Consecutive Executions framework I

ALICE

Helga

Introduction

Problem
space

Subspace of
instances

Feature space

Algorithm
space

Performance
space

Footprints in
instance space

Preference set

Preference
learning

Conclusions

The thesis introduced a framework for learning (linear) composite priority dispatching rule – using job-shop as a case-study – with the following guidelines:

- ★ For a given problem domain, use a suitable problem generator to train and test on.
- ★ Define features to grasp the essence of visited k -solutions
- ★ **Success** is highly dependent on the preference pairs introduced to the system:
 - ★ Ψ_p reduces the preference set without loss of performance.
 - ★ **Stepwise bias** is needed to balance time dependent Ψ_p in order to create **time independent** models.

It is non intuitive how to go about collecting training data.



Using Analysis & Learning Iterative Consecutive Executions framework I

ALICE

Helga

Introduction

Problem space

Subspace of instances

Feature space

Algorithm space

Performance space

Footprints in instance space

Preference set

Preference learning

Conclusions

The thesis introduced a framework for learning (linear) composite priority dispatching rule – using job-shop as a case-study – with the following guidelines:

- ★ For a given problem domain, use a suitable problem generator to train and test on.
- ★ Define features to grasp the essence of visited k -solutions
- ★ **Success** is highly dependent on the preference pairs introduced to the system:
 - ★ Ψ_p reduces the preference set without loss of performance.
 - ★ Stepwise bias is needed to balance time dependent Ψ_p in order to create time independent models.

It is **non intuitive** how to go about **collecting** training data.



Using Analysis & Learning Iterative Consecutive Executions framework II

ALICE

Helga

Introduction

Problem
space

Subspace of
instances

Feature space

Algorithm
space

Performance
space

Footprints in
instance space

Preference set

Preference
learning

Conclusions

Continued from prev. slide:

- ★ Learning **optimal** trajectories predominant in literature. Study showed Φ^{OPT} can result in **insufficient** knowledge.
- ★ Following sub-optimal deterministic policies, yet labelling with an optimal solver, improves the guiding policy.
- ★ In sequential decision making, all future observations are dependent on previous operations. Active update procedure using DAgger ensures sample states the learned model is likely to encounter is integrated to Ψ_p^{DAI} .
- ★ Instead of reusing same problem instances, extend the training set with new instances for quicker convergence of DAgger.



Using Analysis & Learning Iterative Consecutive Executions framework II

ALICE

Helga

Introduction

Problem
space

Subspace of
instances

Feature space

Algorithm
space

Performance
space

Footprints in
instance space

Preference set

Preference
learning

Conclusions

Continued from prev. slide:

- ★ Learning optimal trajectories predominant in literature. Study showed Φ^{OPT} can result in insufficient knowledge.
- ★ Following **sub-optimal** deterministic policies, yet labelling with an optimal solver, **improves** the guiding policy.
- ★ In sequential decision making, all future observations are dependent on previous operations. Active update procedure using DAgger ensures sample states the learned model is likely to encounter is integrated to $\Psi_p^{\text{DA}i}$.
- ★ Instead of reusing same problem instances, extend the training set with new instances for quicker convergence of DAgger.



Using Analysis & Learning Iterative Consecutive Executions framework II

ALICE

Helga

Introduction

Problem space

Subspace of instances

Feature space

Algorithm space

Performance space

Footprints in instance space

Preference set

Preference learning

Conclusions

Continued from prev. slide:

- ★ Learning optimal trajectories predominant in literature. Study showed Φ^{OPT} can result in insufficient knowledge.
- ★ Following sub-optimal deterministic policies, yet labelling with an optimal solver, improves the guiding policy.
- ★ In **sequential** decision making, all future observations are dependent on **previous** operations. Active update procedure using DAgger ensures sample states the learned model is likely to encounter is integrated to $\Psi_p^{\text{DA}i}$.
- ★ Instead of reusing same problem instances, extend the training set with new instances for quicker convergence of DAgger.



Using Analysis & Learning Iterative Consecutive Executions framework II

ALICE

Helga

Introduction

Problem
space

Subspace of
instances

Feature space

Algorithm
space

Performance
space

Footprints in
instance space

Preference set

Preference
learning

Conclusions

Continued from prev. slide:

- ★ Learning optimal trajectories predominant in literature. Study showed Φ^{OPT} can result in insufficient knowledge.
- ★ Following sub-optimal deterministic policies, yet labelling with an optimal solver, improves the guiding policy.
- ★ In sequential decision making, all future observations are dependent on previous operations. Active update procedure using DAgger ensures sample states the **learned model is likely to encounter** is integrated to $\Psi_p^{\text{DA}i}$.
- ★ Instead of reusing same problem instances, extend the training set with new instances for quicker convergence of DAgger.



Using Analysis & Learning Iterative Consecutive Executions framework II

ALICE

Helga

Introduction

Problem
space

Subspace of
instances

Feature space

Algorithm
space

Performance
space

Footprints in
instance space

Preference set

Preference
learning

Conclusions

Continued from prev. slide:

- ★ Learning optimal trajectories predominant in literature. Study showed Φ^{OPT} can result in insufficient knowledge.
- ★ Following sub-optimal deterministic policies, yet labelling with an optimal solver, improves the guiding policy.
- ★ In sequential decision making, all future observations are dependent on previous operations. Active update procedure using DAgger ensures sample states the learned model is likely to encounter is integrated to $\Psi_p^{\text{DA}i}$.
- ★ Instead of reusing same problem instances, extend the training set with **new** instances for **quicker convergence** of DAgger.



Future work

ALICE

Helga

Main conclusions:

★ .

Introduction

Problem
space

Subspace of
instances

Feature space

Algorithm
space

Performance
space

Footprints in
instance space

Preference set

Preference
learning

Conclusions

Thank you for your attention

ALICE

Helga

Introduction

Problem
space

Subspace of
instances

Feature space

Algorithm
space

Performance
space

Footprints in
instance space

Preference set

Preference
learning

Conclusions

Helga Ingimundardóttir
hei2@hi.is

Supplementary material:

- ★ Shiny application
- ★ Github.

