



University of Iceland

Faculty of Industrial Eng., Mechanical Eng. and Computer Science

ALICE

Helga

Introduction

Problem
space

Subspace of
instances

Feature space

Algorithm
space

Performance
space

Footprints in
instance space

Preference set

Preference
learning

Conclusions

ALICE

Analysis & Learning Iterative Consecutive Executions

Helga Ingimundardóttir

University of Iceland

June 30, 2016



Introduction

ALICE

Helga

Introduction

Problem
space

Subspace of
instances

Feature space

Algorithm
space

Performance
space

Footprints in
instance space

Preference set

Preference
learning

Conclusions

Motivation:

- ★ The general goal is to train optimisation algorithms, for an arbitrary problem domain, using data.

Contribution:

- ★ The main contribution of this thesis is towards a better understanding of how this training data should be constructed.

Framework for Algorithm Learning

ALICE

Helga

Introduction

Problem space

Subspace of instances

Feature space

Algorithm space

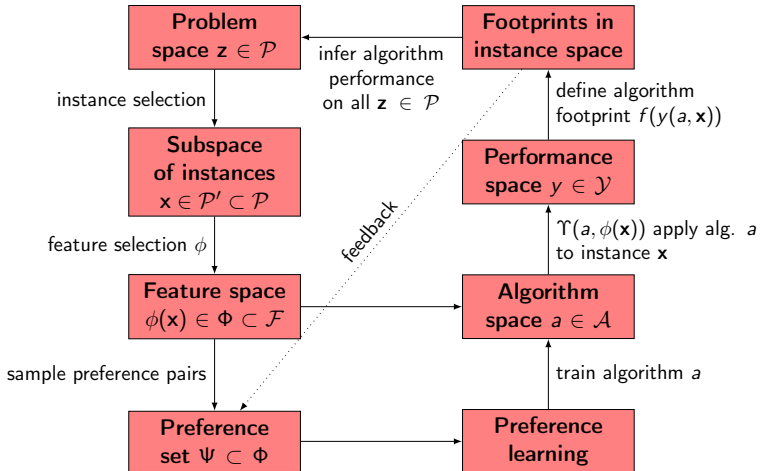
Performance space

Footprints in instance space

Preference set

Preference learning

Conclusions





Mad Hatter tea-party (definition)

ALICE

Helga

Introduction

Problem
space

Subspace of
instances

Feature space

Algorithm
space

Performance
space

Footprints in
instance space

Preference set

Preference
learning

Conclusions

The attending guests: They all have to:

- | | |
|--------------------|---|
| $J_1)$ Alice | $M_1)$ have wine or pour tea |
| $J_2)$ March Hare | $M_2)$ spread butter |
| $J_3)$ Dormouse | $M_3)$ get a haircut |
| $J_4)$ Mad Hatter. | $M_4)$ check the time of the broken watch |
| | $M_5)$ say what they mean. |

This can be considered as is a typical 4×5 job-shop, where:

- ★ our guests are the **jobs**
- ★ their tasks are the **machines**
- ★ objective is to **minimise** C_{\max} (when Alice can leave).

Mad Hatter tea-party (k -solutions)

ALICE

Helga

Midway: $k = 10$




Figure: Disjunctive graph




Figure: Gantt chart

Mad Hatter tea-party (K -solutions)

ALICE

Helga

Introduction

Problem space

Subspace of instances

Feature space

Algorithm space


Performance space

Footprints in instance space

Preference set

Preference learning

Conclusions



Problem instance generators

ALICE

Helga

Introduction

Problem
space

Subspace of
instances

Feature space

Algorithm
space

Performance
space

Footprints in
instance space

Preference set

Preference
learning

Conclusions

	name	size ($n \times m$)	N_{train}	N_{test}	note
JSP	$\mathcal{P}_{j.\text{rnd}}^{6 \times 5}$	6×5	500	500	random
	$\mathcal{P}_{j.\text{rndn}}^{6 \times 5}$	6×5	500	500	random-narrow
	$\mathcal{P}_{j.\text{rnd}, J_1}^{6 \times 5}$	6×5	500	500	random with job variation
	$\mathcal{P}_{j.\text{rnd}, M_1}^{6 \times 5}$	6×5	500	500	random with machine variation
	$\mathcal{P}_{j.\text{rnd}}^{10 \times 10}$	10×10	300	200	random
	$\mathcal{P}_{j.\text{rndn}}^{10 \times 10}$	10×10	300	200	random-narrow
	$\mathcal{P}_{j.\text{rnd}, J_1}^{10 \times 10}$	10×10	300	200	random with job variation
	$\mathcal{P}_{j.\text{rnd}, M_1}^{10 \times 10}$	10×10	300	200	random with machine variation
	$\mathcal{P}_{\text{JSP.ORLIB}}$	various	–	82	various
FSP	$\mathcal{P}_{f.\text{rnd}}^{6 \times 5}$	6×5	500	500	random
	$\mathcal{P}_{f.\text{rndn}}^{6 \times 5}$	6×5	500	500	random-narrow
	$\mathcal{P}_{f.\text{jc}}^{6 \times 5}$	6×5	500	500	job-correlated
	$\mathcal{P}_{f.\text{mc}}^{6 \times 5}$	6×5	500	500	machine-correlated
	$\mathcal{P}_{f.\text{mxc}}^{6 \times 5}$	6×5	500	500	mixed-correlation
	$\mathcal{P}_{f.\text{rnd}}^{10 \times 10}$	10×10	300	200	random
	$\mathcal{P}_{\text{FPS.ORLIB}}$	various	–	31	various



Features for JSP

ALICE

Helga

Introduction

Problem
space

Subspace of
instances

Feature space

Algorithm
space

Performance
space

Footprints in
instance space

Preference set

Preference
learning

Conclusions

job	ϕ_1	job processing time
	ϕ_2	job start-time
	ϕ_3	job end-time
	ϕ_4	job arrival time
	ϕ_5	time job had to wait
	ϕ_6	total processing time for job
	ϕ_7	total work remaining for job
	ϕ_8	number of assigned operations for job
machine	ϕ_9	when machine is next free
	ϕ_{10}	total processing time for machine
	ϕ_{11}	total work remaining for machine
	ϕ_{12}	number of assigned operations for machine
	ϕ_{13}	change in idle time by assignment
	ϕ_{14}	total idle time for machine
	ϕ_{15}	total idle time for all machines
	ϕ_{16}	current makespan
final makespan	ϕ_{17}	final makespan using SPT
	ϕ_{18}	final makespan using LPT
	ϕ_{19}	final makespan using LWR
	ϕ_{20}	final makespan using MWR
	ϕ_{RND}	final makespans using 100 random rollouts
	ϕ_{21}	mean for ϕ_{RND}
	ϕ_{22}	standard deviation for ϕ_{RND}
	ϕ_{23}	minimum value for ϕ_{RND}
	ϕ_{24}	maximum value for ϕ_{RND}

Trajectory strategies for ϕ

ALICE

Helga

Introduction

Problem
space

Subspace of
instances

Feature space

Algorithm
space

Performance
space

Footprints in
instance space

Preference set

Preference
learning

Conclusions

Following the **policy**:

- ★ (ϕ^{OPT}) expert π_* .
- ★ (ϕ^{SPT}) shortest processing time (SPT).
- ★ (ϕ^{LPT}) longest processing time (LPT).
- ★ (ϕ^{LWR}) least work remaining (LWR).
- ★ (ϕ^{MWR}) most work remaining (MWR).
- ★ (ϕ^{RND}) random policy (RND).
- ★ $(\phi^{\text{ES}.\rho})$ the policy obtained by optimising with CMA-ES.
- ★ (ϕ^{ALL}) union of all of the above.

Sampled size of $|\Phi(k)|$

$(6 \times 5, N_{train} = 500)$

ALICE

Helga

Introduction

Problem space

Subspace of instances

Feature space

Algorithm space


Performance space

Footprints in instance space

Preference set

Preference learning

Conclusions





Various Methods for Solving JSP

ALICE

Helga

Introduction

Problem space

Subspace of instances

Feature space

Algorithm space

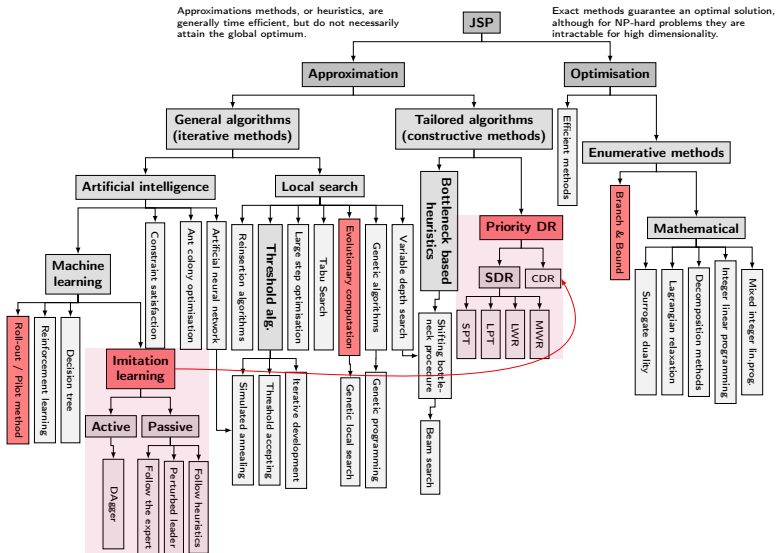
Performance space

Footprints in instance space

Preference set

Preference learning

Conclusions





Performance measure

ALICE

Helga

Introduction

Problem
space

Subspace of
instances

Feature space

Algorithm
space

Performance
space

Footprints in
instance space

Preference set

Preference
learning

Conclusions

Performance of policy π compared with its optimal makespan, found using an expert policy, π_* , is the following loss function:

$$\rho = \frac{C_{\max}^{\pi} - C_{\max}^{\pi_*}}{C_{\max}^{\pi_*}} \cdot 100\%$$

The goal is to minimise this discrepancy between **predicted** value and **true** outcome.

Deviation from optimality, ρ

ALICE

Helga

Introduction

Problem space

Subspace of instances

Feature space

Algorithm space



Performance space

Footprints in instance space

Preference set

Preference learning

Conclusions



Simple dispatching rule

- Shortest Processing Time
- Longest Processing Time
- Least Work Remaining
- Most Work Remaining
- Random dispatches

Data set

- train
- test

Making optimal decisions, ξ_π^*

ALICE

Helga

Introduction

Problem
space

Subspace of
instances

Feature space

Algorithm
space

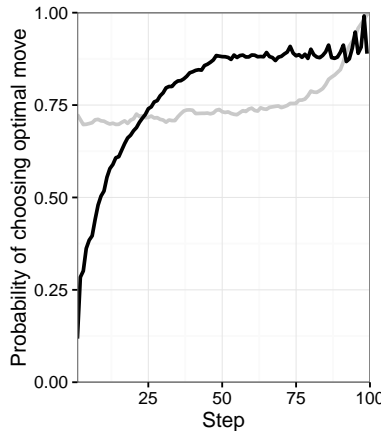

Performance
space

Footprints in
instance space

Preference set

Preference
learning

Conclusions



Probability of SDR being optimal, $\xi_{\langle \text{SDR} \rangle}^*$

ALICE

Helga

Introduction

Problem
space

Subspace of
instances

Feature space

Algorithm
space


Performance
space

Footprints in
instance space

Preference set

Preference
learning

Conclusions



Blended dispatching rule

ALICE

Helga

Introduction

Problem
space

Subspace of
instances

Feature space

Algorithm
space


Performance
space

Footprints in
instance space

Preference set

Preference
learning

Conclusions



Impact of suboptimal decision, $\{\zeta_{\min}^*, \zeta_{\max}^*\}$

ALICE

Helga

Introduction

Problem
space

Subspace of
instances

Feature space

Algorithm
space


Performance
space

Footprints in
instance space

Preference set

Preference
learning

Conclusions



Probability of SDR being optimal, $\xi_{\langle \text{SDR} \rangle}$

ALICE

Helga

Introduction

Problem
space

Subspace of
instances

Feature space

Algorithm
space


Performance
space

Footprints in
instance space

Preference set

Preference
learning

Conclusions



Impact of suboptimal decision, $\{\zeta_{\min}^{\pi}, \zeta_{\max}^{\pi}\}$

ALICE

Helga

Introduction

Problem
space

Subspace of
instances

Feature space

Algorithm
space

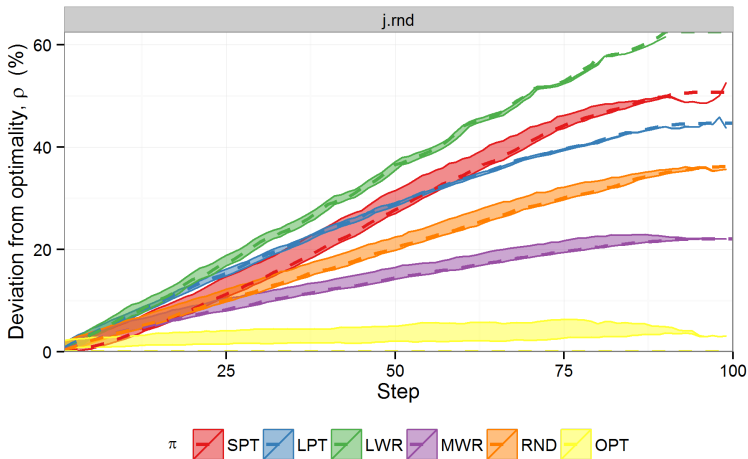
Performance
space

Footprints in
instance space

Preference set

Preference
learning

Conclusions





Generating training data

ALICE

Helga

Introduction

Problem
space

Subspace of
instances

Feature space

Algorithm
space

Performance
space

Footprints in
instance space

Preference set

Preference
learning

Conclusions

ALICE framework for creating dispatching rules:

- ★ **Linear classification** to identify good dispatches, from worse ones.
- ★ **Generate** feature set, $\Phi \subset \mathcal{F}$, both from
 - ★ **optimal** solutions, ϕ^o
 - ★ **suboptimal** solutions, ϕ^sby exploring various **trajectories** within the feature-space (where $\phi^o, \phi^s \in \mathcal{F}$).
- ★ Sample Φ to **create** training set Ψ with rank pairs:
 - ★ **optimal** decision, $(z^o, y_o) = (\phi^o - \phi^s, +1)$
 - ★ **suboptimal** decision, $(z^s, y_s) = (\phi^s - \phi^o, -1)$using different **ranking** schemes (where $z^o, z^s \in \Psi$)
- ★ **Sample** Ψ using **stepwise bias** for time independent policy.

Sampled size of $|\Psi(k)|$ (6×5 , $N_{train} = 500$)

ALICE

Helga

Introduction

Problem
space

Subspace of
instances

Feature space

Algorithm
space

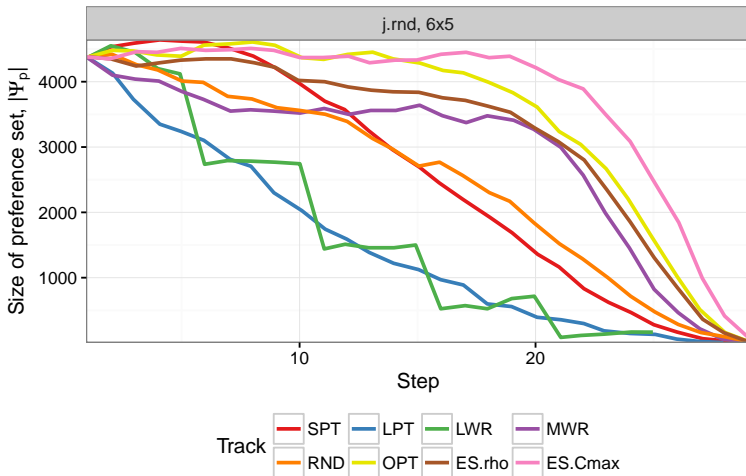
Performance
space

Footprints in
instance space

Preference set

Preference
learning

Conclusions



Stepwise bias strategies

ALICE

Helga

Introduction

Problem
space

Subspace of
instances

Feature space

Algorithm
space


Performance
space

Footprints in
instance space

Preference set

Preference
learning

Conclusions





Ordinal Regression

ALICE

Helga

Introduction

Problem
space

Subspace of
instances

Feature space

Algorithm
space

Performance
space

Footprints in
instance space

Preference set

Preference
learning

Conclusions

Preference learning:

- ★ Mapping of points to ranks: $\{h(\cdot) : \Phi \mapsto Y\}$ where

$$\phi_o \succ \phi_s \iff h(\phi_o) > h(\phi_s)$$

- ★ The preference is defined by a **linear** function:

$$h(\phi) = \langle \mathbf{w}_i \cdot \phi \rangle$$

optimised w.r.t. \mathbf{w} based on training data Ψ

- ★ Note: **Limitations** in **approximation** function to capture the complex dynamics incorporated in optimal trajectories.

Various Methods for Solving JSP

ALICE

Helga

Introduction

Problem space

Subspace of instances

Feature space

Algorithm space

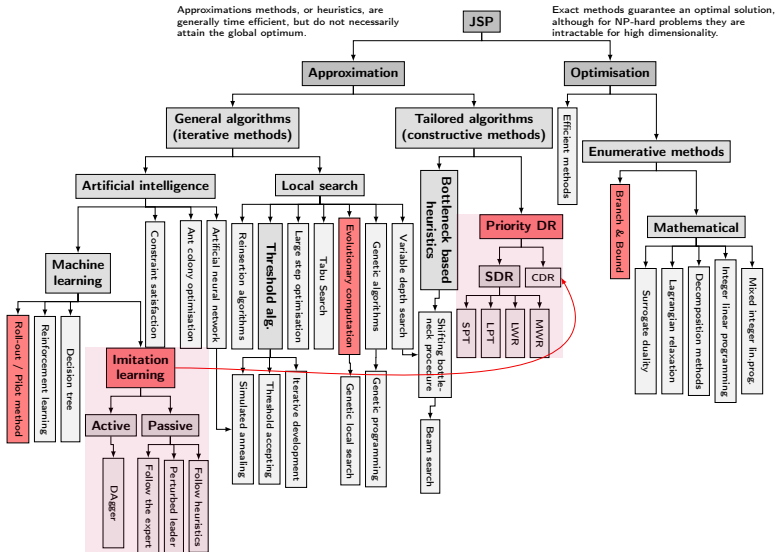
Performance space

Footprints in instance space

Preference set

Preference learning

Conclusions



Passive imitation learning

ALICE

Helga

Introduction

Problem
space

Subspace of
instances

Feature space

Algorithm
space

Performance
space

Footprints in
instance space


Preference set

Preference
learning

Conclusions

Passive imitation learning (single pass):

- ★ Prediction with expert advice, π_*
- ★ Follow the perturbed leader (OPT_ϵ)
- ★ Follow a heuristic (e.g. SDRs).



Active imitation learning

ALICE

Helga

Introduction

Problem
space

Subspace of
instances

Feature space

Algorithm
space

Performance
space

Footprints in
instance space

Preference set

Preference
learning


Conclusions

Active imitation learning (**iterative**):

★ Dataset Aggregation (**DAgger**)

$$\pi_i = \beta_i \pi_{\star} + (1 - \beta_i) \hat{\pi}_{i-1}$$

where $\hat{\pi}_{i-1}$ is the previous learned model, and $\hat{\pi}_i$ learns on **aggregated dataset** of all previous iterations.



Deviation from optimality, ρ

ALICE

Helga

Introduction

Problem space

Subspace of instances

Feature space

Algorithm space

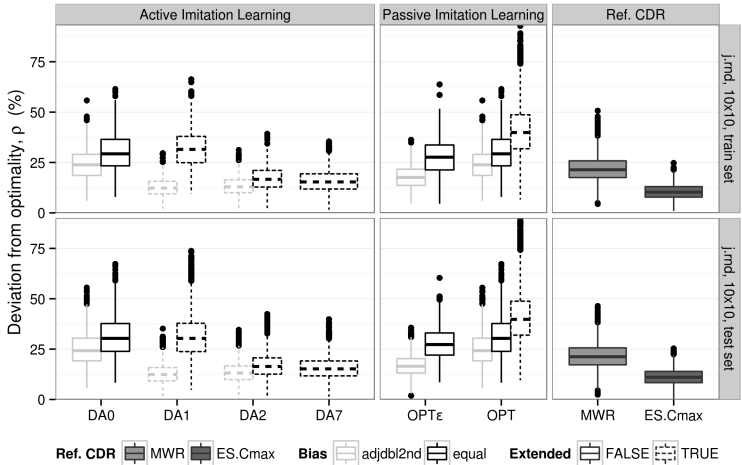
Performance space

Footprints in instance space

Preference set

Preference learning

Conclusions





Using Analysis & Learning Iterative Consecutive Executions framework I

ALICE

Helga

Introduction

Problem space

Subspace of instances

Feature space

Algorithm space

Performance space

Footprints in instance space

Preference set

Preference learning

Conclusions

The thesis introduced a framework for learning (**linear**) composite priority dispatching rule – using **job-shop** as a case-study – with the following guidelines:

- ★ For a given problem domain, use a suitable problem **generator** to **train** and **test** on.
 - ★ Define **features** to grasp the essence of visited **k -solutions**
 - ★ **Success** is highly dependent on the preference pairs introduced to the system:
 - ★ Ψ_p **reduces** the preference set **without loss** of performance.
 - ★ **Stepwise bias** is needed to balance time dependent Ψ_p in order to create **time independent** models.
- It is **non intuitive** how to go about **collecting** training data.



Using Analysis & Learning Iterative Consecutive Executions framework II

ALICE

Helga

Introduction

Problem
space

Subspace of
instances

Feature space

Algorithm
space

Performance
space

Footprints in
instance space

Preference set

Preference
learning

Conclusions

Continued from prev. slide:

- ★ Learning **optimal** trajectories predominant in literature. Study showed Φ^{OPT} can result in **insufficient** knowledge.
- ★ Following **sub-optimal** deterministic policies, yet labelling with an optimal solver, **improves** the guiding policy.
- ★ Active update procedure using **DAgger** ensures sample states the **learned model is likely to encounter** is integrated to $\Psi_p^{\text{DA}i}$.
- ★ Instead of reusing same problem instances, extend the training set with **new** instances for **quicker convergence** of DAgger.
- ★ In **sequential** decision making, all future observations are dependent on **previous** operations.

Acknowledgements

ALICE

Helga

Introduction

Problem
space

Subspace of
instances

Feature space

Algorithm
space

Performance
space

Footprints in
instance space

Preference set

Preference
learning

Conclusions

Funding: University of Iceland's
Research Fund.

Doctoral committee:

- ★ Prof. Tómas Philip Rúnarsson,
University of Iceland (advisor).
- ★ Prof. Gunnar Stefánsson,
University of Iceland.
- ★ Prof. Michèle Sebag,
Université Paris-Sud.



Thank you for your attention

ALICE

Helga

Introduction

Problem
space

Subspace of
instances

Feature space

Algorithm
space

Performance
space

Footprints in
instance space

Preference set

Preference
learning

Conclusions

Questions?

Helga Ingimundardóttir
hei2@hi.is

Supplementary material:

- ★ Shiny application
- ★ Github.

